



**Barcelona
Supercomputing
Center**
Centro Nacional de Supercomputación

OpenMP Tasking Introduction

Supercomputing and Parallel Programming

Xavier Teruel (BSC) and Xavier Martorell (BSC)



Online, April 22nd, 2024

Introduction to OpenMP Tasking



OpenMP traditionally supports the fork-join model

- With worksharing primitives: for, sections, single

Now we will introduce tasking in a way that integrates perfectly with the fork-join model

With tasking,

- Parallel regions are created with the PARALLEL directive, as usual
- Immediately inside the PARALLEL, we use SINGLE, in order to restrict the task creation to a single thread
 - » The single thread encounters TASKS and creates them
 - » When reaching the end of SINGLE, it proceeds to execute tasks
 - » Any additional threads are also dedicated to execute tasks

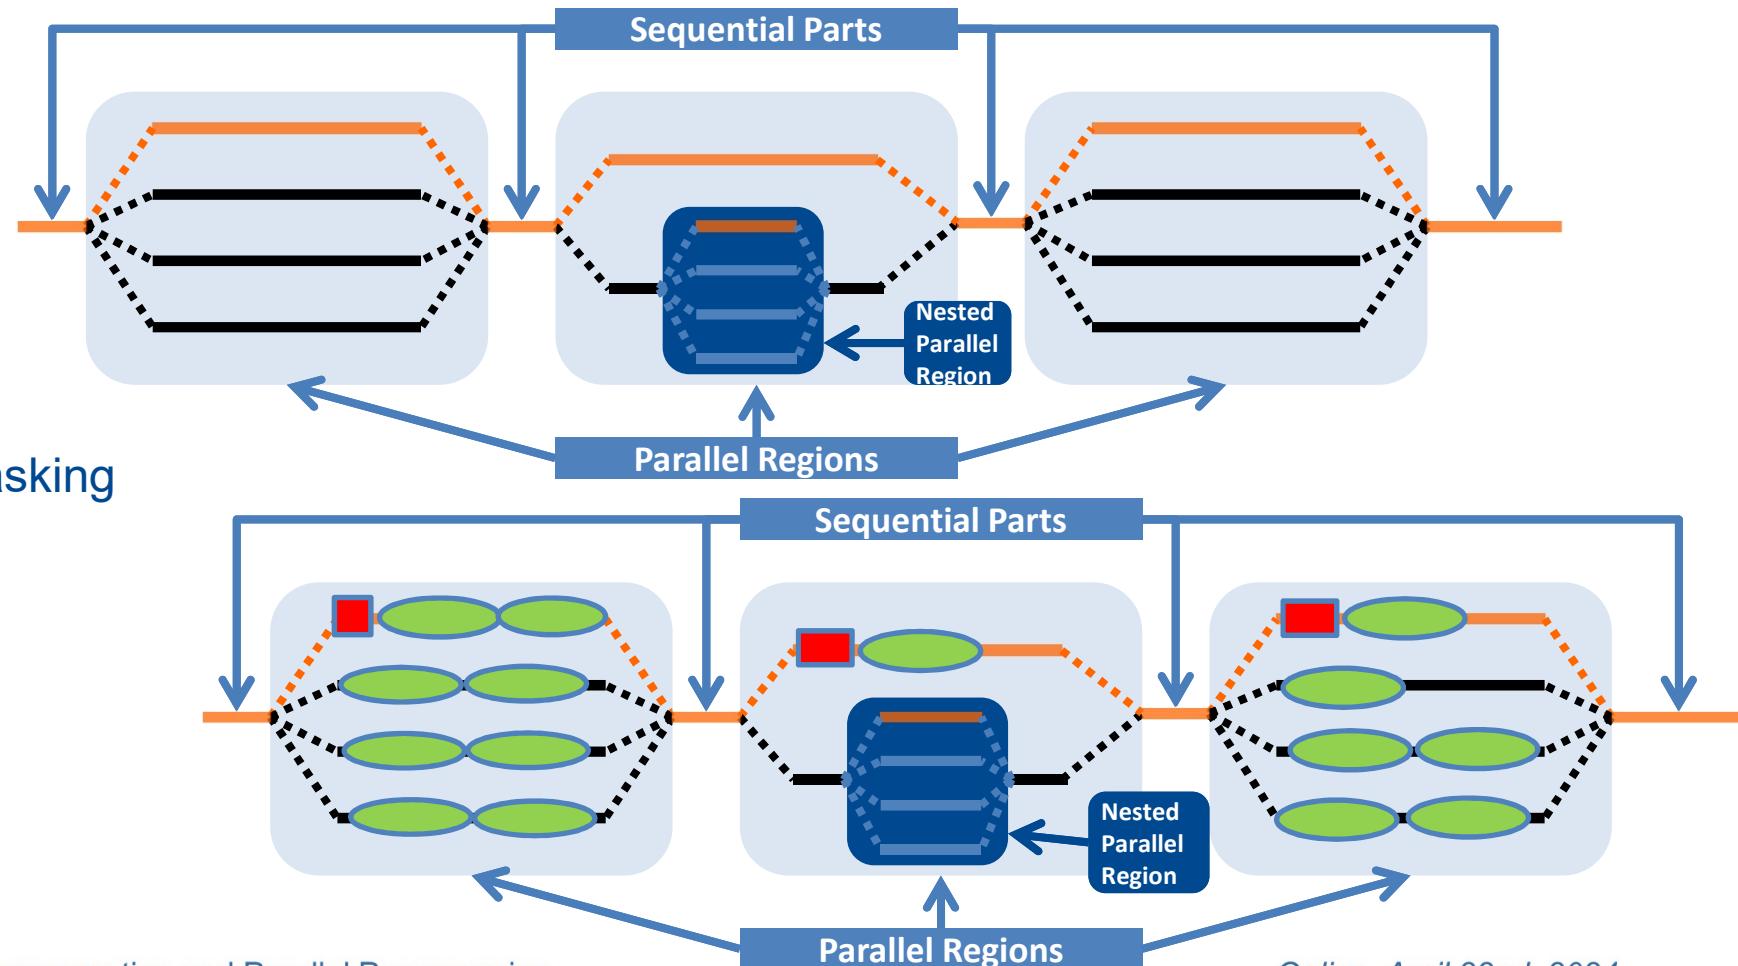
Data environment

- Very similar to the fork-join approach
 - » Task variables are FIRSTPRIVATE by default
 - » ... can also be used as PRIVATE or SHARED

Reminding the fork-join model



Traditional fork-join



www.bsc.es



**Barcelona
Supercomputing
Center**
Centro Nacional de Supercomputación

OpenMP Tasking Basics

Supercomputing and Parallel Programming



Online, April 22nd, 2024

What is a task in OpenMP?



Tasks are **work units** whose execution may be deferred...
... or it can be executed immediately!!!

Tasks appears in OpenMP 3.0 specification (2008)

Tasks are composed of:

- code to execute (set of instructions, function calls, etc...)
- a data environment (initialized at creation time)
- internal control variables (ICVs)

In OpenMP tasks are created...

- when reaching a parallel region → implicit task are created per thread
- when encounters a task construct → explicit task is created
- when encounters a taskloop construct → explicit task per chunk is created
- when encounters a target construct → target task is created

Tasking execution model

Supports unstructured parallelism

- unbounded loops

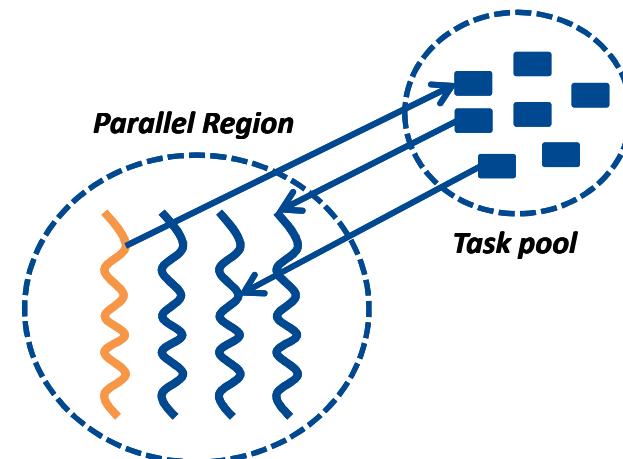
```
while ( <expr> ) {  
    ...;  
}
```

- recursive function calls

```
void myCode ( <args> ) {  
    ...; myCode ( <args> ); ...;  
}
```

Several scenarios are possible

- single creator vs. multiple creators...
- but all members in the team are candidates to execute these tasks



The task construct

Deferring a unit of work (executable for any member of the team)

- always attached to a structured block

```
#pragma omp task [clause[,] clause]...
{structured-block}
```

Where clause:

- private(list), firstprivate(list), shared(list)
- default(shared | none)
- untied
- if(scalar-expression)
- mergeable
- final(scalar-expression)
- priority(priority-value)
- depend(dependence-type: list)

Task data environment: what is the default?



Pre-determined data-sharing attributes

- threadprivate variables are threadprivate
- dynamic storage duration objects are shared (malloc, new,...)
- static data members are shared
- variables declared inside the construct
 - » static storage duration variables are shared
 - » automatic storage duration variables are private
- the loop iteration variable(s)...

Explicit data-sharing clauses (shared, private, firstprivate,...)

- if default clause present, what the clause says
 - » none means that the compiler will issue an error if the attribute is not explicitly set by the programmer (very useful!!!)

Implicit data-sharing rules for the task region

- the shared attribute is lexically inherited
- in any other case the variable is firstprivate

Task default data-sharing attributes (in practice)



```
int a ;
void foo ( int b ) {
    int c;
    #pragma omp parallel private( c )
{
    int d ;
    #pragma omp task
{
    int e;
    a = <expr>;
    b = <expr>;
    c = <expr>;
    d = <expr>;
    e = <expr>;
    g = <expr>;
}
}
}
```

- `default(None)` may help when you are not sure of understand the default

Task scheduling: tied vs untied tasks (1)



Tasks are tied by default (when no untied clause present)

- tied tasks are executed always by the same thread (*not necessarily creator*)
- tied tasks “may” run into performance problems

Programmers may specify tasks to be untied (relax scheduling)

```
#pragma omp task untied  
{structured-block}
```

- can potentially switch to any thread (of the team)
- bad mix with thread based features: thread-id, threadprivate, critical regions...
- gives the runtime more flexibility to schedule tasks

Task scheduling: tied vs untied tasks (2)

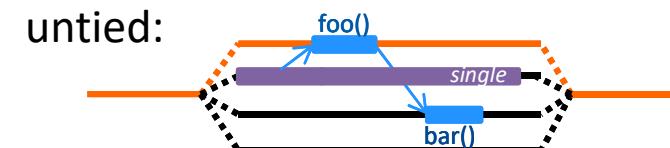
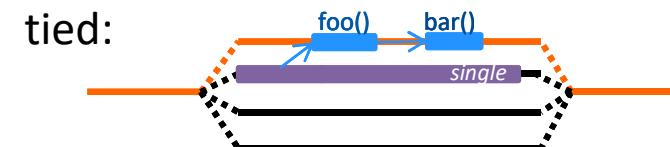
Task scheduling points (and the taskyield directive)

- tasks can be suspended/resumed at these points
- some additional constraints to avoid deadlock problems
- implicit scheduling points (creation, synchronization, ...)
- explicit scheduling point: the taskyield directive

```
#pragma omp taskyield
```

Scheduling untied tasks: example

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task [untied]
    {
        foo ();
        #pragma omp taskyield
        bar ();
    }
}
```



Controlling task scheduling (1)

The **if** clause of a task construct

- allows to optimize task creation/execution → reduces parallelism but also reduces the pressure in the runtime's task pool
- for “very” fine grain tasks you may need to do your own (manual) if

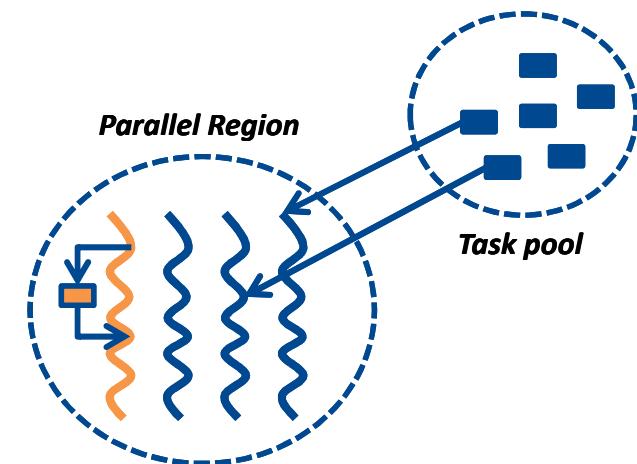
```
#pragma omp task if(expresion)  
{structured-block}
```

If the expression of the “if” clause evaluates to false

- the encountering task is suspended
- the new task is executed immediately
- the parent task resumes when the task finishes

This is known as **undispatched task**

...more combined with mergeable clause!!!



Controlling task scheduling (2)



The **mergeable** clause of a task construct

- allows to optimize task creation/execution (combined with the if clause)
- under certain circumstances it may avoid the whole task overhead

```
#pragma omp task mergeable [if(expression)]  
{structured-block}
```

if-clause evaluates to false → task is executed immediately

- But with its own data environment and ICVs

Combined with the semantic of the mergeable clause

- “a task for which the data environment (inclusive of ICVs) may be the same as that of its generating task region”
- so the user agrees (if possible) on relaxing the previous restriction

Undeferrable and mergeable task may execute as a function call

But it will never be possible when there are private variables

Controlling task scheduling (3)



The **final** clause of a task construct

- allows to omit future task creation → reduces parallelism & overhead

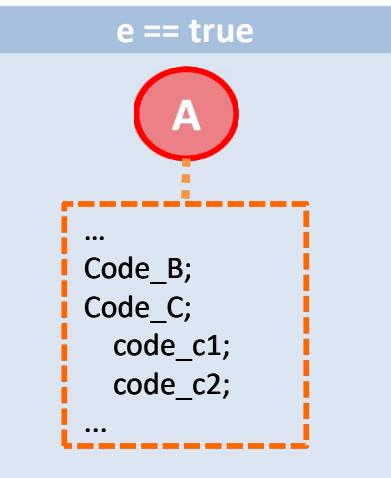
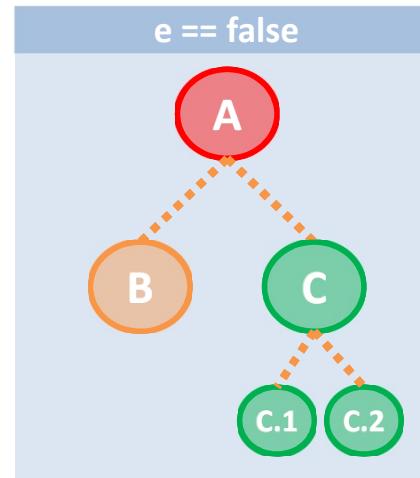
```
#pragma omp task final(expresion)
{structured-block}
```

If the expression of the “final” clause evaluates to true

- the new task is created and executed normally
- in the context of this task no new tasks will be created

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task final(e)
    {
        #pragma omp task
        { code_B; }
        #pragma omp task
        { code_C; }
        #pragma omp taskwait
    }
}
```

Children tasks may have additional task constructs



Programmer's hints for task scheduler

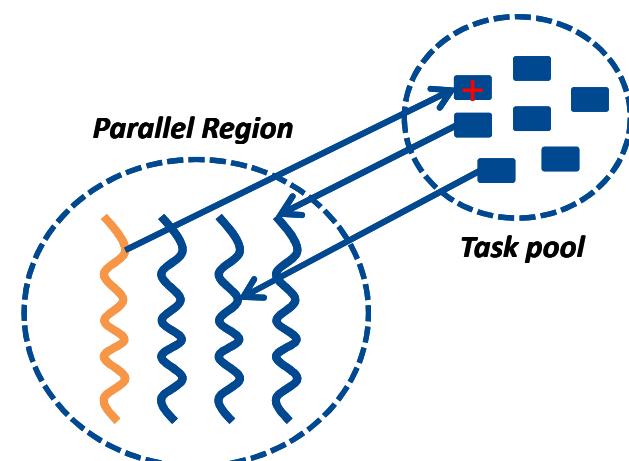


Programmers may specify a priority value when creating a task

```
#pragma omp task priority(pvalue)
{structured-block: loop}
```

- pvalue: the higher → the best (will be scheduled earlier)
- all ready tasks are inserted in an ordered ready queue
- once a thread becomes idle, gets one of the highest priority tasks

```
#pragma omp parallel
#pragma omp single
{
    for ( i = 0; i < SIZE; i++) {
        #pragma omp task priority(1)
        { code_A; }
    }
    #pragma omp task priority(100)
    { code_C; }
    ...
}
```





www.bsc.es



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

Task Synchronization

Supercomputing and Parallel Programming

Online, April 22nd, 2024

Synchronizing the execution of tasks



Threads need “some” order in the sequence of their actions

- execute in a logical order certain regions
- mutual exclusion in the execution of a given region
- wait in a location until all other threads have reach the same location
- wait until a given condition is accomplished

OpenMP provides different synchronization mechanisms

- master construct → already explained in previous sessions
- critical construct → already explained in previous sessions
- barrier directive → already explained, but...
- atomic construct → already explained in previous sessions
- taskwait directive
- taskgroup construct
- depend clause

The barrier directive (and tasks)

Threads cannot proceed until all threads have reach the barrier
... and all previously generated work is completed!!!

```
#pragma omp barrier
```

- Some constructs have an implicit barrier at the end (e.g., the parallel construct, single, sections, {for/do} loop,...)

Using barrier to force task completion

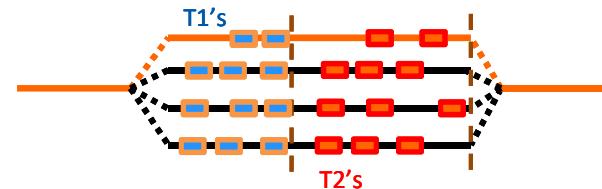
```
#pragma omp parallel
{
    #pragma omp master
    generate_tasks_T1 ();

    #pragma omp barrier

    #pragma omp master
    generate_tasks_T2 ();
}
```

Forces all tasks (T1) to be executed

Implicit barrier: also forces tasks to complete



Waiting for child tasks

The taskwait directive (shallow task synchronization)

- It is a stand-alone directive

```
#pragma omp taskwait
```

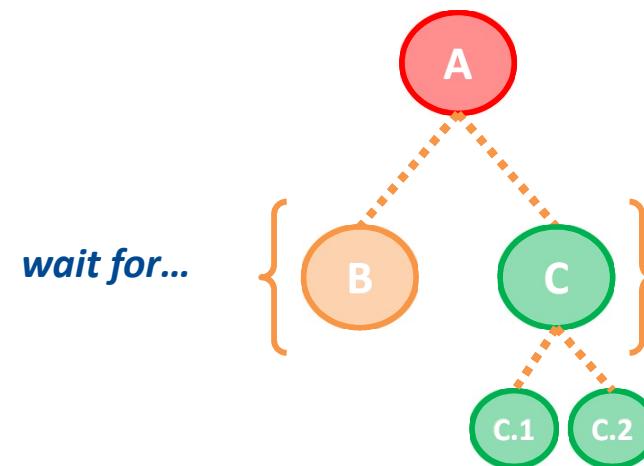
- wait on the completion of child tasks of the current task
- just direct children, not descendants
- includes an implicit task scheduling point

Using the taskwait directive

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    {
        #pragma omp task
        { ... }
        #pragma omp task
        { ... }
    }
    #pragma omp taskwait
}
```

Children tasks may create additional tasks

Wait only for direct descendant tasks



Waiting for all descendant tasks

The taskgroup construct (deep task synchronization)

- always attached to a structured block

```
#pragma omp taskgroup  
{structured-block}
```

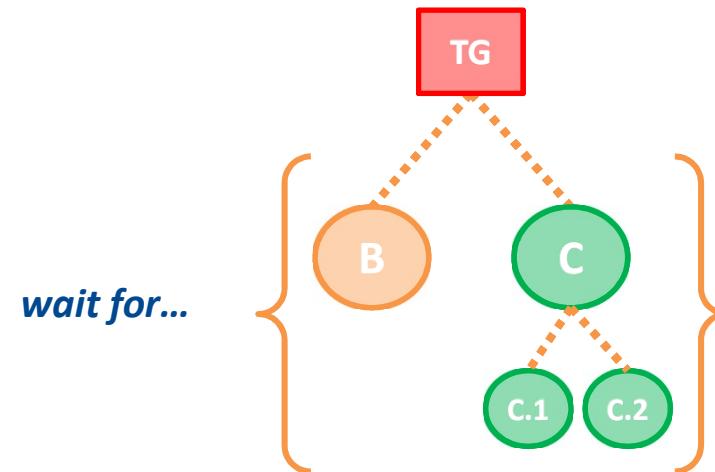
- wait on the completion of all descendant tasks of the current task
- includes an implicit task scheduling point at the end of the construct

Using the taskgroup construct

```
#pragma omp parallel  
#pragma omp single  
{  
    #pragma omp taskgroup  
    {  
        #pragma omp task  
        [...]  
        #pragma omp task  
        [...]  
    }  
}
```

Children tasks may create additional tasks

Wait for all descendant tasks



www.bsc.es



**Barcelona
Supercomputing
Center**
Centro Nacional de Supercomputación

OpenMP Tasking Intermediate

Supercomputing and Parallel Programming



Online, April 22nd, 2024

Using task dependences

The depend clause of the task construct

```
#pragma omp task depend(dependence-type: list)
{structured-block}
```

- used to compute dependences, but actually it is not a dependence
- specify the data directionality of a list of variables

Where dependence-type can be:

- in: the task only reads from the data specified
- out: the task only writes to the data specified
- inout: the task reads from and writes to the data

And where list items are

- variables, a named data storage block (memory address)
- array sections, a designated subset of the elements of an array
 - » A[lower:length]

Computing task dependences (1)

If a task does “in” on a given data variable

- the task will depend on all previously generated sibling tasks that reference at least one of the list items in an out or inout dependence list

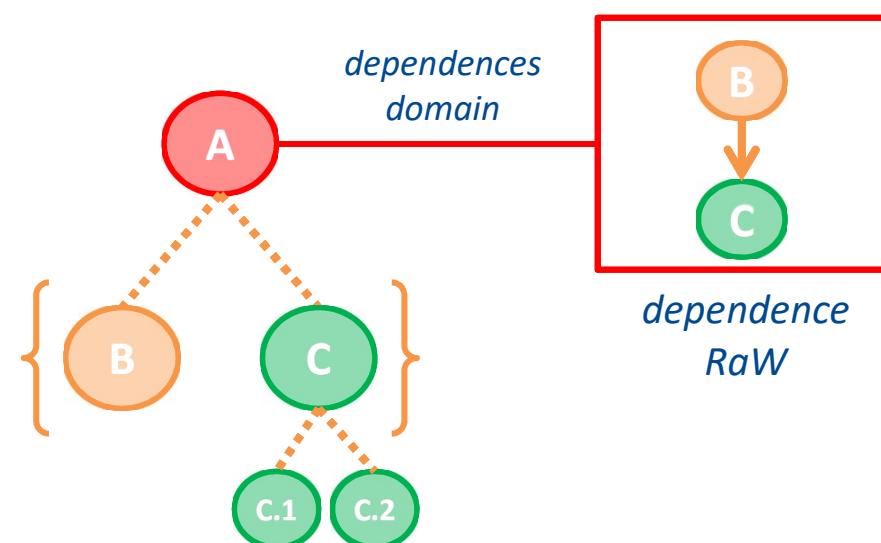
If a task does “out” or “inout” on a given data variable

- on both out and inout dependence types, the task will depend on all previously generated sibling tasks that reference at least one of list items in an in, out or inout dependence list

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    {
        #pragma omp task depend(out:a)
        { ... }

        #pragma omp task depend(in:a)
        { ... }

        #pragma omp taskwait
    }
}
```



Computing task dependences (2)



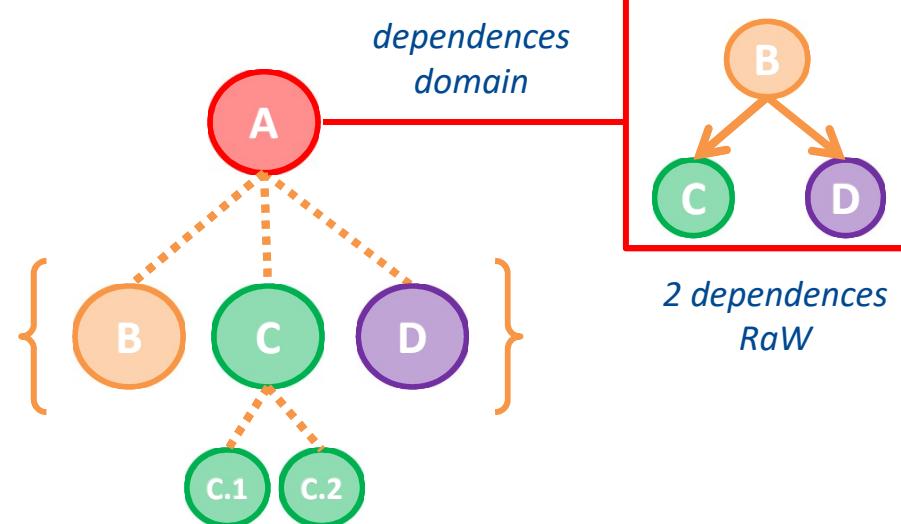
Computing dependences between one writer and n-readers

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    {
        #pragma omp task depend(out:a)
        { ... }

        #pragma omp task depend(in:a)
        { ... }

        #pragma omp task depend(in:a)
        { ... }

        #pragma omp taskwait
    }
}
```



Computing task dependences (3)



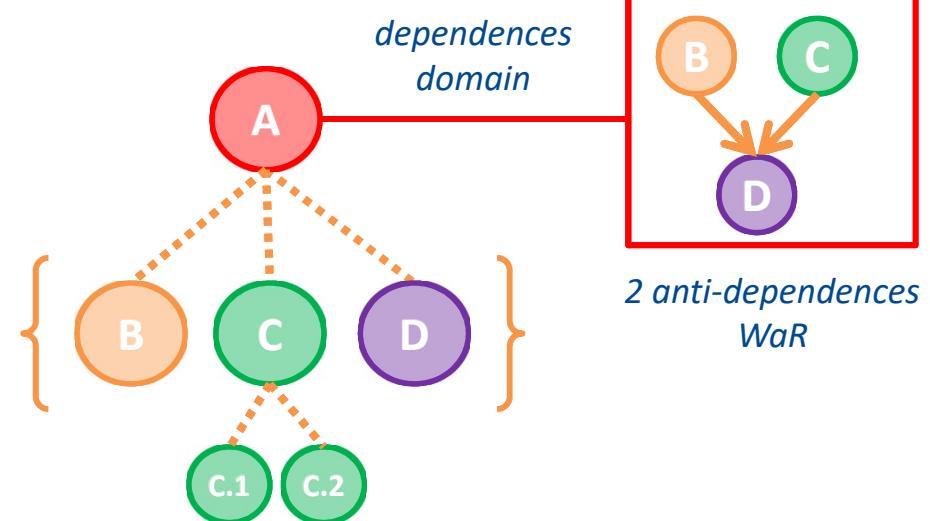
Computing dependences between n-readers and one writer

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    {
        #pragma omp task depend(in:a)
        { ... }

        #pragma omp task depend(in:a)
        { ... }

        #pragma omp task depend(out:a)
        { ... }

        #pragma omp taskwait
    }
}
```

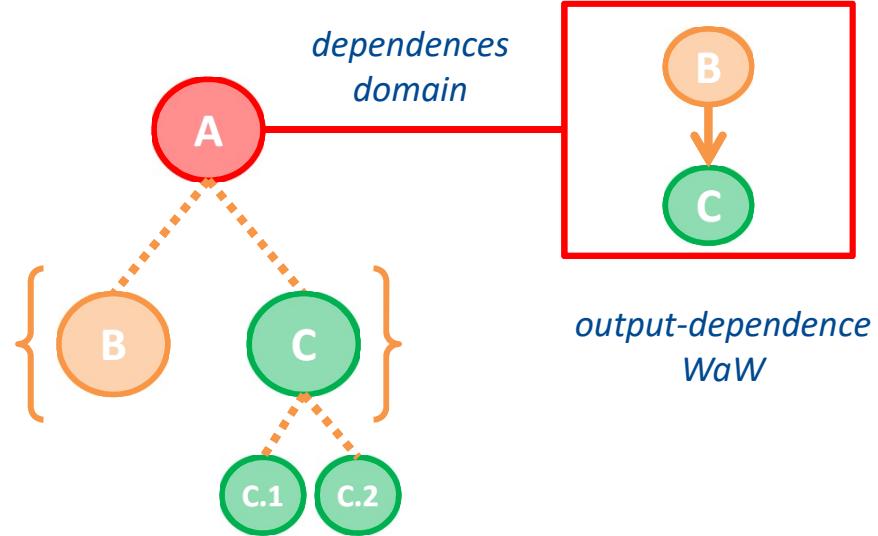


Computing task dependences (4)



Computing dependences between 2 writers

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    {
        #pragma omp task depend(out:a)
        { ... }
        #pragma omp task depend(out:a)
        { ... }
        #pragma omp taskwait
    }
}
```



Using task dependences (cont.)

The depend clause of the task construct

```
#pragma omp task depend(dependence-type: list)
{structured-block}
```

Restrictions on list items

- list items used in depend clauses of the same task or sibling tasks must indicate **identical storage or disjoint storage**
- list items used in depend clauses cannot be zero-length array sections
- a variable that is part of another variable (such as a field of a structure) but is not an array element or an array section cannot appear in a depend clause

```
#define N 100

#pragma omp task depend(out: a[0:N])
{ ... }

#pragma omp task depend(in: a[25:50])
{ ... }
```

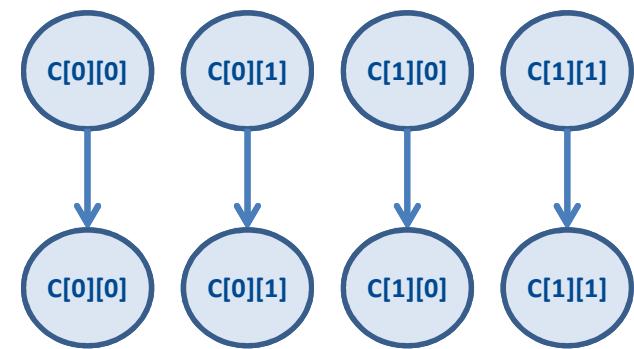


Example: matrix multiply (dependences)



```
void matmul_block ( int N, int BS, float *A, float *B, float *C) ;  
  
// Assume BS divides N perfectly  
void matmul ( int N, int BS, float A[N][N], float B[N][N], float C[N][N] )  
{  
    #pragma omp parallel  
    #pragma omp single  
    {  
        int i, j, k;  
        for ( i = 0; i < N; i+=BS) {  
            for ( j = 0; j < N; j+=BS) {  
                for ( k = 0; k < N; k+=BS) {  
                    #pragma omp task depend ( in:A[i:BS][k:BS],B[k:BS][j:BS] )\  
                        depend ( inout:C[i:BS][j:BS] )  
                    matmul_block (N, BS, &A[i][k], &B[k][j], &C[i][j] );  
                }  
            }  
        }  
    }  
}
```

- avoid “blocks” to be written before read
- input deps useless in this particular example (*still recommended*)
- example on a matrix of 2x2 blocks:





www.bsc.es



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

Task Loop

Supercomputing and Parallel Programming

Online, April 22nd, 2024

Task loop: motivation

Loop (worksharing) construct restrictions

- all threads (in the current team) must reach the worksharing construct
- taskloop constructs comes to break this specific restriction (using tasks)

So if we are executing a single or a section...

```
#include "synthetic.h"

void main (void)
{
    #pragma omp parallel
    #pragma omp sections
    {
        #pragma omp section
        synthetic_phase1();
        #pragma omp section
        synthetic_phase2();
        #pragma omp section
        synthetic_phase3();
    }
}
```

```
#include "synthetic.h"

void synthetic_phase2()
{
    #pragma omp for
    for ( i = 0; i < N ; i ++ ) { ... }
}
```

A large red 'X' icon is positioned to the right of the code snippet, indicating that the shown code is incorrect or invalid.

```
#include "synthetic.h"

void synthetic_phase2()
{
    #pragma omp taskloop
    for ( i = 0; i < N ; i ++ ) { ... }
}
```

A large green checkmark icon is positioned to the right of the code snippet, indicating that the shown code is correct or valid.

The taskloop construct

Deferring several units of work (exec. for any team member)

- always attached to a “for” loop (“do” in Fortran)

```
#pragma omp taskloop [clause[,] clause]...]  
{structured-block: loop}
```

Where clause:

- if(scalar-expr) → already explained (applies to each created task)
- shared(list), private(list), firstprivate(list), lastprivate(list) and default(dtype)
- grainsize(grain-size) and num_tasks(num-tasks)
- collapse(n)
- final(scalar-expr) → already explained (applies to each created task)
- priority(priority-value) → already explained (applies to each created task)
- untied → already explained (applies to each created task)
- mergeable → already explained (applies to each created task)
- nogroup

Using grainsize in taskloop construct

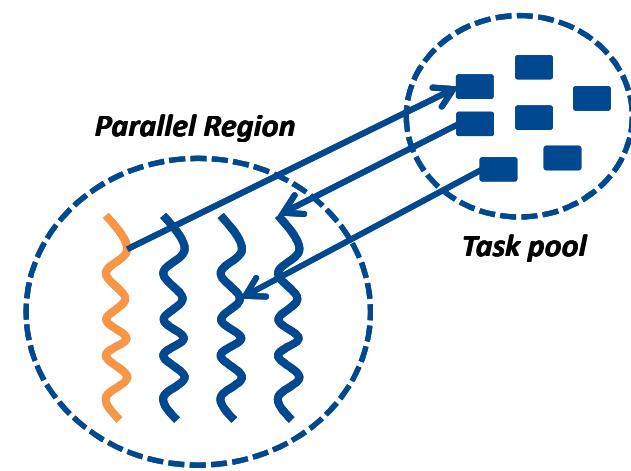
The grainsize clause of the taskloop construct

```
#pragma omp taskloop grainsize(<grain-size>)
{structured-block: loop}
```

- allow to specify the grain size of the generated chunks (tasks)
 - » greater or equal than min(grain-size, iters)
 - » less than two times grain-size ($2 \times$ grain-size)
- cannot be combined with num_tasks clause

```
#include "synthetic.h"

void synthetic_phase2() {
    #pragma omp taskloop grainsize(10)
    for ( i = 0; i < N ; i ++ ) { ... }
}
```



Philosophy: amount of work that is worthy to execute as a task

Using num_tasks in taskloop construct

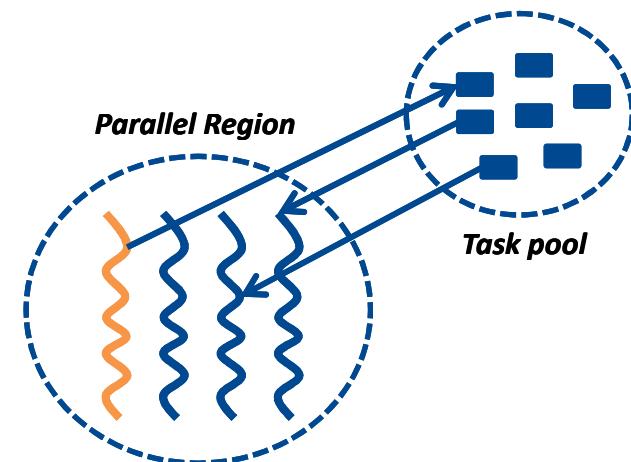
The num_tasks clause of the taskloop construct

```
#pragma omp taskloop num_tasks(<num-tasks>)
{structured-block: loop}
```

- allow to specify the number of chunks (tasks)
 - » greater or equal than min(num-tasks, iters)
 - » each task should have as minimum one iteration
- cannot be combined with the grainsize clause

```
#include "synthetic.h"

void synthetic_phase2() {
    #pragma omp taskloop num_tasks(10)
    for ( i = 0; i < N ; i ++ ) { ... }
}
```



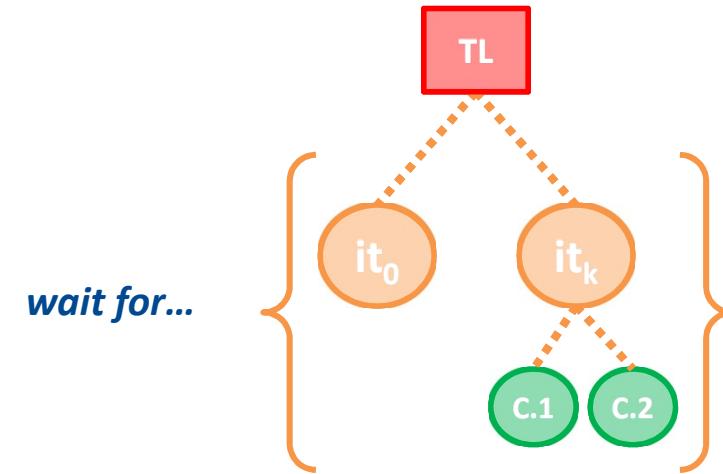
Philosophy: amount of parallelism we want to create

Taskgroup associated with a taskloop



```
#include "synthetic.h"
void synthetic_phase2()
{
    #pragma omp taskgroup
    {
        #pragma omp taskloop
        for ( i = 0; i < N ; i ++ ){ ... }
    }
    foo();
    bar();
}
```

```
#include "synthetic.h"
void synthetic_phase2()
{
    {
        #pragma omp taskloop nogroup
        for ( i = 0; i < N ; i ++ ){ ... }
    }
    foo();
    bar();
}
```



The nogroup clause of the taskloop construct

```
#pragma omp taskloop nogroup
{structured-block: loop}
```

- allow to continue the execution of the encountering task without waiting for all created tasks



www.bsc.es



**Barcelona
Supercomputing
Center**
Centro Nacional de Supercomputación

Thank you!

For further information please visit/contact

xavier.martorell@bsc.es

Intellectual Property Rights Notice: The User may only download, make and retain a copy of the materials for his/her use for non-commercial and research purposes. The User may not commercially use the material, unless has been granted prior written consent by the Licensor to do so; and cannot remove, obscure or modify copyright notices, text acknowledging or other means of identification or disclaimers as they appear. For further details, please contact BSC-CNS.

Supercomputing and Parallel Programming

Online, April 22nd, 2024